

Mohamed Elsayed

University Avenue, Edmonton, AB

✉ mohamedelsayed@ualberta.ca 🌐 mohmdelsayed.github.io

Research Objective.....

Developing general-purpose mechanisms for continual deep reinforcement learning agents, particularly resource-constrained embodied agents that learn in deployment

Education

Doctoral of Philosophy

University of Alberta, Canada

Computing Science. Advisor: A. Rupam Mahmood

GPA: 4.0/4.0

Expected Graduation: Sep 2026

Master of Science

University of Alberta, Canada

Computing Science. Advisor: A. Rupam Mahmood

GPA: 3.9/4.0

2019-2022

Bachelor of Science

University of science and technology, Zewail City, Egypt

Major: Communication and Information Engineering. Minor: Physics of the Earth and Universe.

GPA: 3.7/4.0

2014 – 2019

Experience

Huawei

Support Researcher, Intern

Edmonton, Canada

May 2020 - Dec 2020

- Worked on autonomous intersection navigation for the unprotected left turn problem with RL.
- Co-authored a workshop paper in NeurIPS ML4AD.

Nara Institute of Science and Technology

Robotics Research Intern

Nara, Japan

Aug 2018 - Sep 2018

- Worked on Semantic SLAM problem using deep learning techniques.
- Worked on map in-painting techniques.

Valeo

Deep Learning Research Intern

Smart Village, Egypt

Jul 2018 - Aug 2018

- Worked on automotive-related projects (classification, semantic segmentation, and detection).
- Handled new state-of-the-art network architectures and bench-marked on related works.

Teaching

University of Alberta

Graduate Teaching Assistant

Edmonton, Canada

- CMPUT 174: Introduction to the Foundations of Computation
- CMPUT 365: Introduction to Reinforcement Learning
- CMPUT 340: Introduction to Numerical Methods

Sep 2019 - April 2020

Sep 2022 - Dec 2022

Jan 2024 - April 2024

Zewail City of Science and Technology

Junior Teaching Assistant

Cairo, Egypt

Sep 2018 - Dec 2018

- CIE 417: Machine Learning

Publications

Conferences and Journals:

- Vasan, G., **Elsayed, M.**, Azimi, S. A., He, J., Shahriar, F., Bellinger, C., White, W., & Mahmood, A. R. (2024). Deep Policy Gradient Methods Without Batch Updates, Target Networks, or Replay Buffers. *Neural Information Processing Systems (NeurIPS)*.
- **Elsayed, M.**, Lan, Q., Lyle, C., & Mahmood, A. R. (2024). Weight Clipping for Deep Continual and Reinforcement Learning. *Reinforcement Learning Journal*, 5(1), 2198–2217.
- **Elsayed, M.**, Farrahi, H., Dangel F., & Mahmood, A. R. (2024). Revisiting Scalable Hessian Diagonal Approximations for Applications in Reinforcement Learning. *International Conference on Machine Learning (ICML)*.
- **Elsayed, M.**, & Mahmood, A. R. (2024). Addressing loss of plasticity and catastrophic forgetting in continual learning. *International Conference on Learning Representations (ICLR)*.
- Zhou, M., Luo, J., Vilella, J., Yang, Y., Rusu, D., Miao, J., ..., **Elsayed M.**, ..., & Wang, J. (2021). SMARTS: Scalable Multi-Agent Reinforcement Learning Training School for Autonomous Driving. *Conference on robot learning (CoRL)*.
- Elgammal, M. A., Elkhoully, O. A., Elhosary, H., **Elsayed, M.**, Mohieldin, A. N., Salama, K. N., & Mostafa, H. (2018). Linear and nonlinear feature extraction for neural seizure detection. *International Midwest Symposium on Circuits and Systems (MWSCAS)*.

Pre-prints and Workshops:

- **Elsayed, M.**, Vasan, G., & Mahmood, A. R. (2024). Streaming Deep Reinforcement Learning Finally Works. *arXiv preprint arXiv:2410.14606*.
- **Elsayed, M.**, Vasan, G., & Mahmood, A. R. (2024). Deep Reinforcement Learning Without Experience Replay, Target Networks, or Batch Updates. *NeurIPS Workshop on Fine-Tuning in Modern Machine Learning*.
- **Elsayed, M.**, Mahmood, A. R. (2022) Utility-based Perturbed Gradient Descent: An Optimizer for Continual Learning. *NeurIPS Workshop on Optimization for Machine Learning*.
- **Elsayed, M.**, Mahmood, A. R. (2022) HesScale: Scalable Computation of Hessian Diagonals. *NeurIPS Workshop on Higher-Order Optimization in Machine Learning*.
- **Elsayed, M.**, Hassanzadeh, K., Nguyen, N. M., Alban, M., Zhu, X., Graves, D., & Luo, J. (2020). ULTRA: A reinforcement learning generalization benchmark for autonomous driving. *NeurIPS Workshop on Machine Learning for Autonomous Driving*.

Talks

Streaming Deep Reinforcement Learning is Now Possible <i>Openmind Research Institute Retreat</i>	Oct 2024
Utility-based Representation Search <i>Tea Time Talk, University of Alberta</i>	Aug 2023
Towards Continual Learning Optimizers <i>RL Sofa, Mila</i>	April 2023
Utility-based Perturbed Gradient Descent <i>Lifelong Reinforcement Learning Barbados Workshop</i>	February 2023

Service

- **Reviewer** at ICLR 2023, CoLLAs 2024, RLC 2024, NeurIPS 2024, ICLR 2025, AISTATS 2025.
- **President** of IEEE Zewail City Student Branch (2017) (Granted University Best Leader Award).

Technical Skills

- **Programming Languages:** Python, C/C++
- **Software Libraries:** Pytorch, Jax
- **Technologies:** Git, Bash, $\text{\LaTeX}$